CS221 Vision Project Report

Team: Cognitive Dissonance - Alec Go, Elizabeth Lingg, Anand Madhavan

March 20, 2009

1 Description of the data structures and code

For the most part, our code is organized such that if there is a class, there is usually a header (.h) and source file (.cpp) associated with it by that name. In cases where it is not, the class may have been combined into a single file and in such cases, we explicitly mention the file locations. The descriptions were organized based on various functionalities.
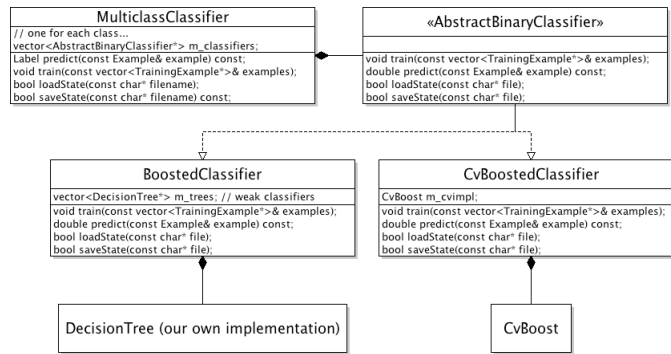
1.1 Classifier

Figure 1: Class diagram for our multi-class classifier

Our classifier code uses a bit of object hierarchy and interface mechanisms to enable quick testing with different implementations. This is described below:

classifier.h/.cpp: Modified code from initial template. Contains the class *AbstractMulticlassClassifier*, which currently has only one implementation in *MulticlassClassifier* described below. The idea behind the *AbstractMulticlassClassifier* interface was to be able to extend this to SVMs and other classifiers at a later stage if required.

MulticlassClassifier.h: This is our main multiclass classifier. This class contains one *AbstractBinaryClassifier* per class (eg: one for mug, one for keyboard etc). The underlying implementation of the *AbstractBinaryClassifier* interface can be one of either *BoostedClassifier* or *CvBoostedClassifier* (which uses *CvBoost* under the hood). This is represented as a class diagram in Figure 1. The depth of a tree and the number of trees are specifyable through constructor arguments to *MulticlassClassifier*.

BoostedClassifier.h: This is our own 'homegrown' implementation based on AdaBoost. It implements the *AbstractBinaryClassifier* interface and further uses class called *DecisionTree* which is our implementation of a simple decision tree that performs binary classification.

CvBoostedClassifier.h: This class also implements the *AbstractBinaryClassifier* interface and uses instead the *CvBoost* implementation of OpenCV.

DecisionTree.h: Our implementation of a simple decision tree that performs binary classification. It uses a n-ary recursive tree structure to store information. It also uses weights for examples, so the sampling distribution can be changed efficiently.

Label.h: Labels for the classifier are specified here. *Label* is currently a typedef to unsigned int. This file also contains functions that map the unsigned int to string and vice versa for efficient usage elsewhere.

NOTE: The performance of our *BoostedClassifier* is comparable to *CvBoost* in precision, recall and other measures (more details in Section 4.2) . However, our final submission version uses *CvBoost* by default, since it is considerably faster than our *BoostedClassifier* implementation.

1.2 Feature Extraction

HaarFeatures.cpp: This class allows us to extract haar features for a given frame.

EdgeDetectionFeatures.cpp: This class allows us to extract sobel (edge detection) features for a given frame.

Hough.cpp: This class allows us to extract hough, canary, and harris (circle, line, edge, and c orner) features for a given frame.

1.3 Motion tracking

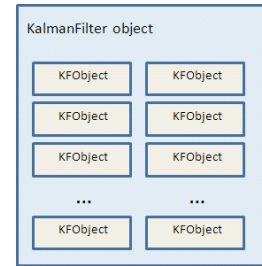Figure 2: Relationship between *LucasKanade* and *LKObject*

Figure 3: Relationship between *KalmanFilter* and *KFObject*

MotionTracker.cpp: This class is a master class that controls which motion tracker will be used (Lucas Kanade or Kalman Filter)

LucasKanade.cpp: This class contains code for Lucas Kanade interpolation and post processing step. The core Lucas Kanade implementation is in LucasKanade.cpp. This file has two main data structures:

- *LucasKanade*: Holds a group of *LKObjects*.

- *LKObject*: Each *LKObject* corresponds to a separate "blob" (held in a CObject) being tracked. This actually contains calls to OpenCV optical flow methods. This relationship is shown in Figure 2.

KalmanFilter.cpp: This class contains code for the Kalman Filter post processing step. The Kalman Filter is used in a similar fashion to the Lucas Kanade motion tracker. The basic data structures are

- *KalmanFilter*: This holds a set of *KFObjects*. The *KalmanFilter* object has an update and predict function, which would trigger all the update and predict processes for the *KFObjects* it holds.

- *KFObject*: This tracks an individual blob. This holds the actual Kalman Filter implementation of (update) and (predict). This relationship is shown in Figure 3.

1.4 Cross validation tool

We introduce a cross validation tool as described in Section 2.1. This can be built using 'make tune'.

tune.cpp*: Main code for doing k-fold cross validation (described in more detail in Section 2.1). Reuses the *CClassifier* code.

Stats.h*: Utility class for storing and computing statistics of precision, recall and F1 scores.

1.5 Infrastructure code

Other files contain infrastructure utilities for rapid development:

Timer.h: A utility class for printing out timing information for functions.

CommandOptions.h: A utility for processing command line arguments.

FinalSettings.h: This is a one stop to override all the command line options since we want *test* and *train* executables to work with the final fixed values of these settings for final submission (and not rely on command line inputs).

2 Implemented Extensions

We implemented a number of useful extensions. Some of these were just diagnosis tools and classifiers, others were features for specific objects and some others were for motion detection.

2.1 k-Fold cross validation tool

We implement a utility to perform k-fold cross validation of our features and classifier. The utility (called 'tune') takes various command line options, such as size of the tree to be used and depth of tree to be used and performs a k-fold cross validation using examples chosen at random.

We go through all the training examples and shuffle them up. This gives us a uniform distribution of samples. We then use the first 'n' examples (specified by a command line argument) and perform k-fold cross validation on it. The first 'fold' is chosen as test set and trained on the rest of the folds. Then the second fold is chosen as the test set and trained on the rest of the folds and so on for all folds. We report the average test as well as the average training error. We report the precision and recall for each category. Finally we also report the confusion matrix. Many other parameters can be specified as command line options as well, these are documented in the Appendix 6.1.

2.2 AdaBoost decision tree implementation

We implement a fully functional AdaBoost[3] based boosted decision tree (as class *BoostedClassifier*). This class uses exponential weighting of falsely predicted examples and the weights are passed along to the nodes in a simple decision tree classifier. The tolerance is determined based on the average feature values encountered. We present results of using our classifier in Section 4.2. We note that although this classifier performs quite well, it is quite slow in comparison to the CvBoost implementation. For the final submission we choose instead to use the CvBoost implementation for its superior speed.

2.3 Shape based features

Based on the type of objects we are given to classify (eg: clocks, keyboards etc), we extract relevant shape features using circle, line, edge and corner detection algorithms.

2.4 Histogram of gradients based features

We implement a histogram feature on each patch, with HOG as the theoretical foundation [1]. We describe this in more detail in Section 4.3.4.

2.5 Motion Tracking

We recognize that using information between frames could also improve our accuracy. We have a Kalman Filter and a Lucas Kanade filter. The motion tracker can track several "blobs" at a time. Each blob is a unique object, like a mug.

2.5.1 False Positives

The motion tracker can decrease precision if there are many false positives returned by the classifier. This is because the motion tracker might continue to track a false positive, even if the classifier stopped returning that specific false positive. Coalescing (mentioned in Section 2.6.2) helps clear false positives, which in turn helps the motion tracker. Another simple heuristic we use is to "count" the number of times the classifier returns a specific blob between frames. If the motion tracker stops seeing a blob returned by a classifier over a series of frames, then it discards it.

2.5.2 A Correspondence Problem Between Frames

We implement a method to resolve the correspondence problem when doing motion tracking. For example, the classifier may detect 3 objects in frame 1, and 4 objects in frame 2. We developed a simple heuristic for mapping objects from frame 1 to objects in frame 2. If there is an overlap above a threshold, then we assume they are the same object. For example, if a mug in frame 1 occupies the 50% (the threshold) of the same sp ace as a mug in frame 2, then we can assume they are the same object.

2.5.3 Object Interpolation

We have heuristics for interpolating object results, in the case that the classifier returned an object in frame 1 and frame 3, but didn't find it in frame 2.

In the most trivial case, object interpolation was used as a mechanism for speeding up video processing (running test.cpp). In this case, the decision tree is occassionally skipped and the motion tracker is used to guess the locations of objects. This h ad a very significant impact on speed because our decision tree was several orders of magnitude slower than the motion tracker.

2.6 Optical Flow - Lucas Kanade

We chose to implement the Lucas Kanade algorithm via the cvCalcOpticalFlowPyrLK method. The test videos meet the three assumptions required by this algorithm [2]:

1. Brightness Constancy. We assume that in the test videos, lighting will be consistent.

2. Temporal persistence. The test videos have steady movement. Only the camera moved. The objects in the scene never moved.

3. Spatial coherence

With these three assumptions, the Lucas Kanade algorithm should perform reasonably well with the test videos.

2.6.1 Implementation Details

Lucas Kanade is used for simple interpolation in frame skipping. The classifier.cpp file uses the following procedure:

1. Every time the decision tree is used, it calls LucasKanade.seed(). This allows the LucasKanade object to "remember" which objects it should track.

2. Every time the decision tree is skipped, it calls LucasKanade.interpolate(), which returns a list of CObjects.

LucasKanade is for tracking a set of points, but the decision classifier outputs a set of CObjects. A conversion from a CObject bounding box and a set of points is thus necessary. The conversion process works as follows: LucasKanade uses the cvFindGoodFeaturesToTrack method to find the good corners to track in a pixel. Any corners found outside of the CObject's area are discarded. A new bounding box is then wrapped around the remaining corners, and this becomes the new CObject.

One advantage of this approach is that as objects leave the scene, the bounding box shrinks naturally because the corners being tracked disappear (because they pass the border of the image).

2.6.2 Coalescing overlaping rectangles

We implement coalescing of overlapping rectangles that have the same labels. This is done by looking at the overlap and permitting two rectangles with a certain threshold of overlap to be merged. The rectangles are successively merged until no rectangles can be merged any more.

3 Assumptions

As mentioned in Section 2.6, we assume the videos have brightness constancy, temporal persistence and spatial coherence for optical flow to work.

4 Experimental results

We perform a number of experiments on each of our extensions. We present the results of our various experiments below.

4.1 Baseline classifier

We perform basic analysis of the CvBoost decision tree. We use CvBoost::GENTLE type with a fixed split criteria of 0.95. We use our k-fold cross validation tool to tweak the size of the trees used, as well as the depth of the trees used. We also use our tool to arrive at a 'bang for the buck' set of parameters, so as to minimize development time while still giving good enough accuracies. We use a 4-fold cross validation for all the below experiments. Average test and training errors are reported across these folds. In analyzing our baseline performance, we use only the 57 Haar features from the milestone run.

4.1.1 Effect of maximum depth of tree

Average Errors (10000 examples)

**Time Taken (in seconds)**

Figure 4: Baseline classifier: Effect of increasing depths on error.

Figure 5: Baseline classifier: Time taken with increasing depth of trees.

Using 10000 examples, we notice that varying depths of the tree, we get better accuracies with increasing depths. However we also note that the time taken increases almost disproportionately for the improvements in accuracies obtained (see Figures 4 and 5). For example for depth 1, we see an accuracy of 1.85% accuracy, while for depth 8, this improves to about 1.81%. However given the time it takes (81 seconds vs 419 seconds), we can judge that increasing the depth a practical way of running our development cycle. Thus we arrive at an optimal depth of 1 or 2.

4.1.2 Effect of number of trees used in boosting

Average Error (10000 examples)

**Time Taken (in seconds) (10000 examples)**

Figure 6: Baseline classifier: Effect of number of trees on errors.

Figure 7: Baseline classifier: Time taken with increasing number of trees.

We next perform experiments by varying the number of trees used during boosting in CvBoost. With increasing number of trees we see that we get lower and lower training errors, but also lower and lower test errors upto a point, after which the results seem to plateau (see Figure 6). Doing a time analysis, we notice depths of 1 perform much better on large number of trees (see Figure 7). This makes us pick depth of 1 and use 400 trees for our development. However, note that we use larger number of trees (1000), and depth 2 for our final submission.

### 4.1.3 Effect of number of training examples

Figure 8: Baseline classifier: Effect of number of examples on error

Figure 9: Baseline classifier: Time taken with increasing number of training examples

We perform a basic analysis on the effect of the number of examples on the test error. This gives us insights into how many examples we should use for our development. Here we notice that our curves follow the typical training and testing error curves, with gradually increasing training errors an decreasing test error. Test errors are typically around 1.5% compared to training errors which vary from 0 to 0.3%. At the least this validates our model as not having high bias. Also while there is an obvious benefit from using more data, we also notice that for 16000 examples, we get test errors of around 1.44%. This just takes around 210 seconds, while using all the 46359 examples gives us a test error of 1.33% but takes 699 seconds (an increase of 3 times). From Figure 8, we also notice that the 'knee' of the curve maybe occuring with 16000 examples. We thus use this for development. Of course we use all the data for training our classifier for final submission.

4.2 'Homegrown' Adaboost classifier

Extending our work from the milestone, we also designed a multiclass classifier that implements the AdaBoost algortihm for boosting (in class called *BoostedClassifier*).

	'Homegrown' version	CvBoost version
Average training error	0.00102101	0.00386118
Average test error	0.0169115	0.0133954

Table 1: Comparison of average test and training errors using the two implementations.

The errors compared to the *CvBoost* classifier are shown in Table 1.

'Homegrown' version				CvBoost version			
	Precision	Recall	F1 score		Precision	Recall	F1 score
mug	0.705	0.617	0.658	mug	0.819	0.597	0.690
stapler	0.786	0.672	0.724	stapler	0.828	0.578	0.681
keyboard	0.633	0.638	0.636	keyboard	0.830	0.519	0.639
clock	0.313	0.939	0.469	clock	0.885	0.469	0.613
scissors	0.693	0.779	0.733	scissors	0.900	0.648	0.754
other	0.992	0.991	0.992	other	0.989	0.989	0.994

Table 2: Comparison of precision, recall and F1 scores for the two implementations.

The precision and recall numbers for the two classifiers are shown side-by-side in Table 2. Of course, the two methods use different variations of boosting, so the numbers are only used primarily as a validation of our classifier's performance.

4.3 Object detection features

4.3.1 Hough Transforms (Circle and Line Detection)

Without circle detection				With circle detection			
	Precision	Recall	F1 score		Precision	Recall	F1 score
mug	0.929	0.553	0.693	mug	0.947	0.574	0.715
stapler	0.805	0.522	0.633	stapler	0.793	0.515	0.624
keyboard	0.833	0.489	0.616	keyboard	0.929	0.707	0.802
clock	1.000	0.400	0.571	clock	1.000	0.400	0.571
scissors	0.938	0.571	0.710	scissors	0.953	0.581	0.722
other	0.987	0.998	0.993	other	0.989	0.999	0.994

Table 3: Comparison of precision, recall and F1 scores with and without hough circle detection.

Hough transforms are features that are used for shape detection in order to do object recognition. In this case, we focus on circles and lines within images. Circles are used for detecting clocks and lines are used for detecting all five objects. Before performing the Hough circle transform, we smooth the image with a Gaussian.

We find that circles are a useful feature, allowing us to avoid misclassification and false negatives, most notably for keyboards. We try several circle features such as highlighting circles in the image and perform a histogram on the image based on pixel intensity. We also use the number of circles in the image as a feature. This turns out to be our best feature. The performance improvement using 4 fold cross validation is shown in Table 3.

We find that the Hough line features do not improve performance. This is partially due to background noise in the images as well as the imprecision in the line detection algorithm and the fact that the algorithm produces lines of infinite length (instead of line segments). Some of the features we have considered are the number of lines in the image, as well as highlighting the lines in the image and performing a histogram based on pixel intensity. Future work would be to analyze different line features.

4.3.2 Edge Detection features

Without line detection	Precision	Recall	F1 score	With line detection	Precision	Recall	F1 score
mug	0.929	0.553	0.693	mug	0.929	0.553	0.693
stapler	0.805	0.522	0.633	stapler	0.847	0.537	0.658
keyboard	0.833	0.489	0.616	keyboard	0.947	0.587	0.725
clock	1.000	0.400	0.571	clock	1.000	0.400	0.571
scissors	0.938	0.571	0.710	scissors	0.909	0.571	0.702
other	0.987	0.998	0.993	other	0.988	0.999	0.993

Table 4: Comparison of precision, recall and F1 scores with and without edge detection.

After running Hough transforms, we consider that instead of analyzing the frames for specific shapes, we can use a rough outline of the objects in the image as a feature. Using the Canny edge detector, we are able to create an outline of the object. We then convert the outline from grayscale into a black and white image. The edges of the objects are highlighted in black, with the background in white. We create a histogram feature with two bins, one containing the white pixels, the other containing black pixels.

We find that this was an effective feature overall, improving performance for staplers and keyboards (see Table 4).

4.3.3 Corner Detection features

Without corner detection	Precision	Recall	F1 score	With corner detection	Precision	Recall	F1 score
mug	0.929	0.553	0.693	mug	0.898	0.564	0.693
stapler	0.805	0.522	0.633	stapler	0.796	0.552	0.652
keyboard	0.833	0.489	0.616	keyboard	0.908	0.750	0.821
clock	1.000	0.400	0.571	clock	1.000	0.400	0.571
scissors	0.938	0.571	0.710	scissors	0.969	0.600	0.741
other	0.987	0.998	0.993	other	0.989	0.998	0.994

Table 5: Comparison of precision, recall and F1 scores with and without corner detection.

We guess that our classification of keyboards could benefit from using corner detection algorithms, since there are many corners on the keyboard. We use the Harris Corner Detection algorithm in the cvFindGoodFeatures function to detect the number of corners on an object. The number of corners as such, are used as a feature, and improve the performance for staplers and scissors. Most notably we find a large performance improvement for keyboards, confirming our initial hypothesis. In addition, it is interesting to note that mugs and clocks do not see improvements. This is due to the fact that these objects are curved and do not in have corners. The results are shown in Table 5.

4.3.4 Histogram of Oriented Gradients (HOG)

	Without HOG				With HOG		
	Precision	Recall	F1 score		Precision	Recall	F1 score
mug	0.929	0.553	0.693	mug	0.952	0.638	0.764
stapler	0.805	0.522	0.633	stapler	0.881	0.664	0.757
keyboard	0.833	0.489	0.616	keyboard	0.903	0.707	0.793
clock	1.000	0.400	0.571	clock	1.000	0.400	0.571
scissors	0.938	0.571	0.710	scissors	0.962	0.714	0.820
other	0.987	0.998	0.993	other	0.992	1.000	0.996

Table 6: Comparison of precision, recall and F1 scores with and without HOG.

We use a histogram feature on the gray scale image, with HOG as the theoretical foundation. HOG involves three stages, gradient computation, orientation binning, and block normalization.

First, we implement stage one, by computing the gradient of the image, using a sobel operator. Then we implement stage two, orientation binding, by creating the histogram of the object based on the image intensity values. Finally, we implement stage three, by normalizing the histogram.

After some experimentation, we find that the feature works best on the original gray scale image, without using the sobel operatior. With the sobel operator, we do not achieve a significant performance increase. Further work would include using different operators to compute the gradient of the image.

The results in Table 6 are with histogram normalization and 80 bins (without the sobel operator). We find that we got the same results with or without histogram normalization. Figure 10 shows the effect of number of bins on the F1 scores. We notice that scores improve progressively for clocks with higher bin sizes, while not playing a significant role in the others.
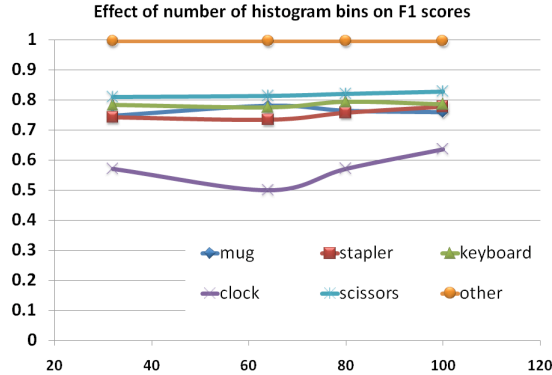

Figure 10: Baseline classifier: Effect of increasing depths on error.

4.4 Effect of coalescing of rectangles during motion tracking

No coalescing	Ovlerap ratio 0.01	Ovlerap ratio 0.2	Ovlerap ratio 0.4	Ovlerap ratio 0.8
0.0438409	0.31638	0.305373	0.27673	0.160496

Table 7: Effect on F1 scores of coalescing overlapping rectangles that are classified as the same object.

We note that when run with the videos, our classifier picks a large number of overlapping patch rectangles as positives. This causes our F1 scores to be dramatically worse than it should be because

false positives are accentuated. We employ an algorithm of coalescing rectangles if they overlap area is more than a certain ratio. Table 7 shows the effect of the overlap ratio on the F1 scores. We coalesce the rectangles if the amount of overlap exceeds the overlap ratio. All the results were obtained using easy.avi.

4.5 Motion tracking experiments

4.5.1 Effect of Lucas Kanade parameters

We perform a few tests to find optimal Lucas Kanade parameters. Because of the dependence on the time dimension, we test these against easy.avi.


Figure 11: Effect of number of frames skipped

Figure 12: Effect of pyramid size

Effect of number of frames skipped:

After a certain number of frames, the decision tree is skipped and the Lucas Kanade interpolation is used instead. Figure 11 shows the number of frames in which the decision tree is skipped, with the Lucas Kanade motion tracker used instead to predict movements. This shows that Lucas Kanade is a reasonable sustitute for the decision tree.

Effect of number of pyramid size:

We vary the number of image pyramids used in Lucas Kanade. A large number of pyramids is useful when there is large global movement (Figure 12). The number of pyramids does not impact performance significantly. This is because the motion in easy.avi is relatively stable. Our Lucas Kanade in the final submission uses 5 pyramids.

Effect of optimal search window size:

We alter the search window size for the optical flow algorithm. This is the size of the window to search for a point from one frame to the next. The results are shown in Figure 13. For easy.avi, the optimal parameter for the window size is 6. Requiring anything higher could potentially waste CPU cycles. Our Lucas Kanade motion tracker in the final submission uses a window size of 10.

Effect of the track error thresholding:

In the optical flow algorithm, we discard corners whose "track error" is too large. The track error is the difference between a patch around a tracked point between two frames. If there is a track error above a certain threshold, Lucas Kanade discards the point. Figure 14 shows how track error affects F1 performance. Our Lucas Kanade motion tracker uses a track error threshold of 300.

Figure 13: Effect of the optical search window size

Figure 14: Effect of the track error thresholding

4.5.2 Kalman Filter

When visualizing the output from the classifier, the CObjects jump around erratically when the classifier processed an image independently of any other image. The Kalman filter can smooth out this erratic movement. Our Kalman Filter uses blobtrackpostprockalman.cpp in cvaux. In this model, the state vector is (x,y,w,h,dx,dy,dw,dh) and the measurement vector is (x,y,w,h).

The major limitation with the Kalman Filter is that it assumes linear dynamics. This assumption held only for the parts of the video where the camera was moving at a constant velocity (the model we chose only had a parameter for velocity). The Kalman Filter improved the F1 results on easy.avi. In one test (with all variables held constant), the Kalman Filter improved the mug F1 score from to 0.414073 to 0.440678. However, when using the Kalman Filter with Lucas Kanade, the results are erratic - sometimes F1 scores increase and sometimes F1 scores decrease. This is because our Kalman Filter's heuristics to remove false postives (mentioned above) are not compatible with Lucas Kanade interpolation.

5 Note on efficiencies and speedups

We use the following to obtain efficiencies during our development process, as well as to obtain efficiencies during the final classification:

- We use integers as labels, instead of strings to save time on map indexing using the labels in many places in the code.

- We use our 'tune' utility (described in Section 2.1) to whet our features on the static images using 4-fold cross validation before using it on the movies.

- We use Lucas Kanade based interpolation techniques (as described in Section 4.5.1) to predict the rectangles in skipped frames to speeden up the testing on movies.

References

[1] N. Dalal and B. Triggs. Histograms of oriented gradients for human detection. volume 1, pages 886–893, 2005.

[2] Bradski G. and Adrian Kaehler. Learning opencv, 2008.

[3] Andrew Ng. Boosting lecture notes cs221, 2009.

6   Appendix

6.1   Tune utility details

6.1.1   Command options

A number of command line options are provided to the tune utility. Some are documented here:

- -examples <integer>: number of examples to use

- -fold <integer> : number of folds to use

- -onefold : boolean that specifies only to do cross-validation on the first of 'k' folds (for quickly testing out new features)

- -depth <integer>: max depth of decision tree to use

- -trees <integer>: number of trees to use for boosting

- -homegrownboost: specifies which version of the classifier is to be used

- -trainerror: option to spit out training error

- -circle_feature, -corner_feature, -edge_feature, -sobel_feature etc: include the various features

6.1.2   Sample output of the tool

Together, these provide us with a valuable tool to diagnose problems before running our classifier on the more time consuming movies. An example of the output of the code is given below:

$ make tune

$ ./tune -homegrownboost -trainerror -fold 4 -examples 16000 -depth 1
-trees 400 -files /afs/ir/class/cs221/vision/data/vision_all
===
Using *HOMEGROWN* boosting classifier
Using 400 trees in boosting
Using 1 depth in boosting
Using 4 folds for validation.
Using files in /afs/ir/class/cs221/vision/data/vision_all
Using 16000 examples
Running experiments on total of: 16000 files
===
Average test errors:
Fold 0: Test error: 0.0185 Training error: 0
Time taken in fold : 166.64 seconds
Fold 1: Test error: 0.0155 Training error: 0
Time taken in fold : 176.69 seconds
Fold 2: Test error: 0.02125 Training error: 0
Time taken in fold : 177.07 seconds
Fold 3: Test error: 0.0175 Training error: 0
Time taken in fold : 175.28 seconds
───────────────────────────────
Avg Training error Avg Test error
0 0.0181875
───────────────────────────────────── -
Confusion Matrix: Predicted labels ->
───────────────────────────────────── -
mug stapl keybo clock sciss other
───────────────────────────────────── -
mug 60 0 2 0 0 32
stapl 3 85 3 0 1 42
keybo 1 0 63 1 2 25
clock 0 0 0 11 0 4
sciss 0 0 3 0 77 25
other 27 32 27 13 48 15413
───────────────────────────────────── -
Prec Recall F-1
───────────────────────────────────── -

13

mug 0.659 0.638 0.649
stapl 0.726 0.634 0.677
keybo 0.643 0.685 0.663
clock 0.440 0.733 0.550
sciss 0.602 0.733 0.661
other 0.992 0.991 0.991

Time taken in entire experiment: 695.68 seconds